

Collaborator or Assistant?

How AI Coding Agents Partition Work Across Pull Request Lifecycles

An agent's work can surface under a human identity; the trace measures visibility, not authorship.

Young Jo Chung · Safwat Hassan · University of Toronto


Alware 2026 · Montreal · 6-7 July

STEP
1

The instrument: PRs as traces

The talk gave the map; the poster uses it to read visibility and review.

1. Each PR is a lifecycle
- A sequence of events, not just a row in a table.
2. The map separates roles
- First commit, review, and merge are distinct trace surfaces.
3. Initiator = first committer
- It records who the commit is logged under — an attribution surface.



The paper's PR-lifecycle state machine (the map).

Why a state machine?


Because the question is not only whether a PR merged. It is what path it took: who appears first, whether review appears, and who executes the merge.

That surface is visible identity, not verified authorship. This is the gap the poster follows.

STEP
2

Figure 3 changes meaning: visibility, not workforce

Same six scenarios, corrected lens: the bars show how agentic work surfaces in the trace.



Their agents act under human identities, using a person's account ID or name.

How to read it

Warm bars: the agent/bot is visible at first commit. Green bars: the first commit surfaces under a human account. In this agentic corpus, green does not mean human-authored; it means agent work is less visible.

Codex carries the pattern (n=20,857, ~99.9% human-signed); Claude is consistent but fragile — n=338 and 57% of its labels flip under attribution-aware re-detection.

STEP
3

Visibility meets review


One association to remember: the least visible agentic work has the least recorded human-approved review.

89.6%

of 29,585 terminal PRs have no recorded human-approved review. Among merged PRs, only 12.8% show one.

17×

spread in human-approved review: Copilot 66.4% to OpenAI Codex 3.9%.



The poster's claim

The lifecycle map still works, but Figure 3 should be read as attribution visibility. Once read that way, the same traces expose a governance gap: human review is sparse overall and lowest where agentic work is least visible as agentic. Association, not cause.

Collaborator or Assistant? A junior teammate — often under your name.

- The puzzle: AI work can enter a repository under a human name.
- The trace: most merges still look human-governed.
- The gap: recorded human-approved review is often absent.
- The question: what should repositories log if they want to govern agent-generated change?



Scan for code & data